

Name these concepts:

"Only one process(/thread) can be in the CS at a time"

"If waiting, then another process can only enter the CS a finite number of times"

"If no other process is in the Critical Section then the process can immediately enter the CS"

Candidate #4

raise my flag if your flag is raised, wait until my turn <i>// Do Critical Section stuff</i> turn = <i>yourid</i> lower my flag	raise my flag if your flag is raised, wait until my turn <i>// Do Critical Section stuff</i> turn = <i>yourid</i> lower my flag
---	---

*// Threads do other stuff and then will repeat in the future*

Problems with 4?

Candidate #5

raise my flag while(your flag is raised) : if it's your turn to win : lower my flag wait while your turn raise my flag <i>// Do Critical Section stuff</i> set your turn to win lower my flag	raise my flag while(your flag is raised) : if it's your turn to win : lower my flag wait while your turn raise my flag <i>// Do Critical Section stuff</i> set your turn to win lower my flag
---	---

#### MYTHS ABOUT THE MUTUAL EXCLUSION PROBLEM

G.L. PETERSON

*Department of Computer Science, University of Rochester, Rochester, NY 14627, U.S.A.*

Received 13 January 1981; revised version received 30 March 1981

Recently in these pages appeared a report by Doran and Thomas [2] which gave partially simplified versions of Dekker-like solutions to the two process mutual exclusion problem with busy-waiting. This report presents a truly simple solution to the problem and attempts in a small way to dispel some myths that seem to have arisen concerning the problem.

Briefly, the *mutual exclusion problem* for two processes is to find sections of code (*trying protocol*, *exit protocol*) for each of two asynchronous processes to use when trying to enter and upon exiting their designated critical sections. The protocols must preserve mutual exclusion and not have deadlock or lockout. *Mutual exclusion* means that both processes can never be in their critical sections at the same time. *No deadlock or lockout* means that no process waits forever inside a protocol. More formal definitions can be found in [5] and elsewhere.

The original solution due to Dekker is discussed at length by Dijkstra in [1]. Of the many reformulations given since, perhaps the best appears in [3]. (Unfortunately the authors believe their correct solution is incorrect.) The solutions of Doran and Thomas are slight improvements which eliminate the 'loop inside a loop' structure of the previously published solutions. The solution presented here has an extremely simple structure and, as shown later, is easy to prove correct.

```
/*trying protocol for P1*/
Q1 := true;
TURN := 1;
wait until not Q2 or TURN = 2;
Critical Section;
/*exit protocol for P1*/
Q1 := false.
```

The protocols of P<sub>1</sub> and P<sub>2</sub> are given in Fig. 1. Q1 and Q2 are initially *false* and TURN may start as either 1 or 2. (The busy wait loop 'wait until Boolean' is just another way of saying "repeat/\* empty statement/\* until Boolean". The Boolean formula is *not* evaluated atomically.)

As can be seen, the algorithm has a very simple structure. This results in an easy proof of correctness. First, neither process can be locked out. Consider P<sub>1</sub>, it has only one wait loop, and assume it can be forced to remain there forever. After a finite amount of time, P<sub>2</sub> will be doing one of three general things: not trying to enter, waiting in its protocol, or repeatedly cycling through its protocols. In the first case, P<sub>1</sub> notes that Q2 is *false* and proceeds. The second case is impossible due to TURN being either 1 or 2; and one of the processes will proceed. In the third case P<sub>2</sub> will quickly set TURN to 2 and never change it back to 1, allowing P<sub>1</sub> to proceed.

If mutual exclusion were not preserved and both processes could somehow end up in their critical sections at the same time, then we have Q1 = Q2 = *true*. Their tests in their wait loops just prior to entering their critical sections at this point could not have been at approximately the same time as TURN would have been favorable to only one of the processes and the other part of the test would have failed for both. This

```
/*trying protocol for P2*/
Q2 := true;
TURN := 2;
wait until not Q1 or TURN = 1;
Critical Section;
/*exit protocol for P2*/
Q2 := false.
```

4 Give two reasons why even implementing a 'correct solution' in C might still fail.

### 3 What are condition variables? How do you use them? What is *Spurious Wakeup*?

Use a condition variable to wait for a data structure to have at least one item. There is one thread that might be calling *pushdata* or delete several times. Another thread that might call *getLast* several times

Plan:

Write a busy wait.

Add mutex

Add condition variable.

<pre>pthread_mutex_t m; pthread_condition_t cv;  float myarray[10]; int count ;  void init() {     pthread_mutex_init(&amp; m, NULL);     pthread_condition_init( &amp; cv, NULL); }  int pushdata(float v) {     myarray[count ++] = v; }  void delete() {     count --; }</pre>	<pre>float getLast() {  // 'result' must be valid! float result = myarray[count];      return result; }</pre>
---	---

### 7. Use a CV to implement a simple version of a *counting semaphore*

Note a real semaphore might implement a queue of waiting threads to ensure fairness (and avoid *starvation*).

<pre>sem_init(sem_t *s, int shared, int value) {  }  </pre>	<pre>typedef struct sem_t {  } sem_t;  </pre>
<pre>sem_post(sem_t*s) {  }  </pre>	<pre>sem_wait(sem_t*s) {  }  </pre>